



*RisingWaveLabs*

# RISING WAVE: THE SQL-NATIVE STREAMING DATABASE

MILESTONE 2  
GROUP 5

## MEMBERS:

ANICA REMIN FERNANDEZ  
ASWIN SIVAKUMAR  
ATHULYA JAYAN  
BETSY VARGHESE  
NAIJEL ROY  
FRIDOLIN FRANKLIN  
SAJIN SAHADEVAN  
ADVAITH HARISH



# agenda

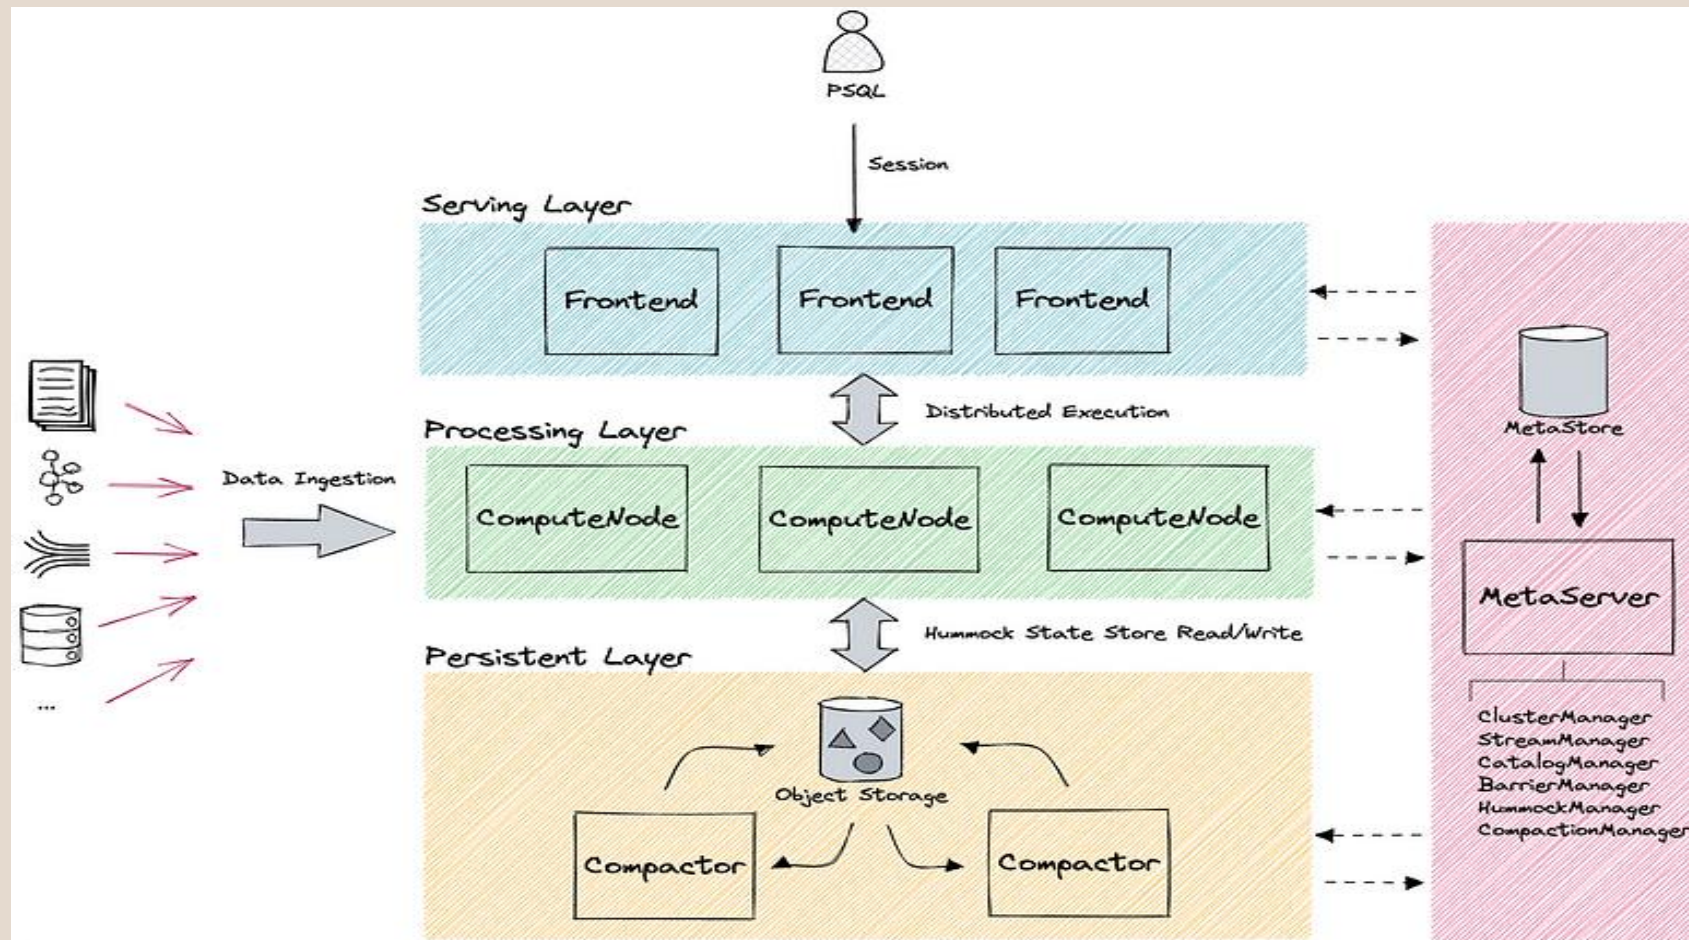
- INTRODUCTION
- CHALLENGES WITH TRADITIONAL STREAMING STACKS
- INTRODUCING RISINGWAVE
- CORE INNOVATIONS
- ADVANCED FEATURES
- REAL-WORLD USE CASES
- LIMITATIONS & CONSIDERATIONS
- CONCLUSION

# What is rising wave ?

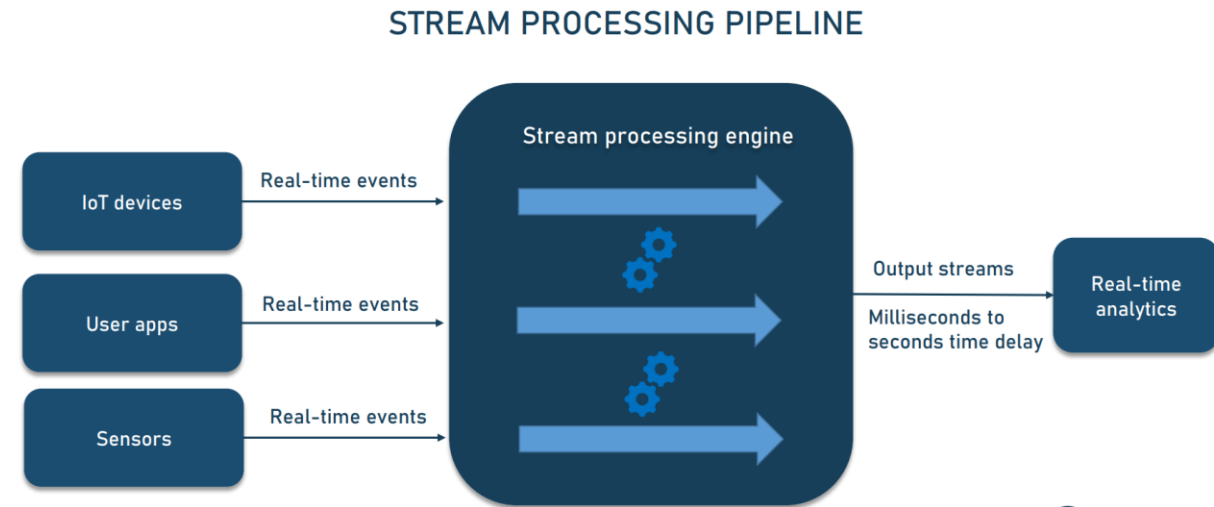
- RISING WAVE IS A SQL NATIVE STREAMING DATABASE BUILT FOR MODERN REAL-TIME ANALYTICS.
- PROVIDES A UNIFIED PLATFORM COMBINING INGESTION, PROCESSING, STORAGE, AND SERVING.
- USES STANDARD POSTGRES SQL WIRE PROTOCOL FOR EASY INTEGRATION. WRITTEN IN RUST FOR PERFORMANCE AND SAFETY.
- ENABLES ELASTIC SCALING BY SEPARATING COMPUTE AND STORAGE LAYERS.
- AIMS TO REDUCE COMPLEXITY AND OPERATIONAL COSTS COMPARED TO TRADITIONAL STACKS.



# RISING WAVE ARCHITECTURE



# What is streaming data?



- It's continuous, real-time data generated from multiple sources. Like IoT devices, social media feeds, financial transactions, and many more.
- It's time-sensitive, requires immediate ingestion and processing.
- Along with batch data its process stored data at intervals.
- Use cases: fraud detection, real-time dashboards, monitoring and alerts.
- Data rates can be thousands to millions per second.

# STREAM PROCESSING VS BATCH PROCESSING

Feature	Batch	Stream
Data Input	Finite datasets	Infinite, real-time flows
Latency	High (minutes–hours)	Low (ms–seconds)
Storage	Data stored, then processed	Processed immediately, minimal storage
Use Cases	Reports, historical analysis	Monitoring, real-time insights
Complexity	Moderate	High (ordering, fault tolerance)
Failure Handling	Re-run entire job	Checkpoints, recovery
Programming Model	SQL, scripts	Event-driven APIs
Examples	ETL, reports	Sensors, fraud detection

# TRADITIONAL STREAMING DATA STACK

- Message Broker :Kafka, Pulsar for data ingestion and buffering.
- Stream Processor: Flink, Spark streaming for computation .
- State Store: RocksDB, Cassandra for maintaining state.
- Challenges include :
  - Complex maintenance across multiple tools.
  - Synchronization and state consistency issues.
  - High operational overhead and specialized skills requirements.
  - Latency introduced by network hops between components.
  - Difficult to implement and scale efficiently.

# Storage vs compute decoupling

Aspect	Traditional Systems	RisingWave
Storage/Compute	Tightly coupled (local storage)	Fully decoupled (object store)
Scalability	Joint scaling needed	Independent scaling
Cost Efficiency	Overprovisioning common	Pay-as-you-grow model
Fault Tolerance	Local state, complex recovery	Persistent snapshots, easy recovery
Resource Use	Less flexible	Dynamic and efficient



# Fault Tolerance & Exactly-One

- Asynchronous checkpointing = minimal query pauses
- Watermarks + snapshots ensure state consistency
- Exactly-once processing
- Resilient to failures with auto recovery
- Enables continuous ops even during scaling or maintenance

## SQL as primary interface

- SQL-first design simplifies stream processing
- Supports:
  - **Materialized Views**
  - **Window Functions**
  - **Temporal Joins**
  - **Event-time semantics**
- Works with PostgreSQL clients, BI tools
- Reduces need for Java/Scala code
- Extensible via **User-Defined Functions (UDFs)**




# Streaming SQL Highlights

- **Continuous Queries:** Instant updates as data flows
- **Windowing:** Tumbling, sliding, session windows
- **Joins:** Stream–stream & stream–table joins
- **Aggregations:** Real-time, low-latency updates
- **Temporal Tables:** Time-travel queries
- Handles **late-arriving data** with watermarking

## Traditional SQL


```
SELECT * FROM YOUR_TABLE WHERE...
```



SQL applied to *static data*

## Streaming SQL

```
SELECT * FROM YOUR_TABLE WHERE...
```



SQL applied to *real-time data stream*

# REAL-TIME USE CASES WITH RISINGWAVE:

## ○ FRAUD DETECTION

- Monitor transactions in real time at scale (millions/sec).
- Detect anomalies using SQL-based pattern matching.
- Instantly flag/block suspicious behavior.
- Enhances compliance and reduces financial loss.

## ○ REAL-TIME AD MONETIZATION

- Count impressions, clicks, and conversions instantly.
- Power dynamic bidding and budget control.
- Feed live dashboards for advertisers/publishers.
- Maximizes revenue with sub-second insights.

## • MARKETING ANALYTICS

- Real-time segmentation with session windows.
- Evaluate A/B tests instantly.
- Live ROI and customer behavior tracking.
- Integrate web, mobile, and social data seamlessly.

## • Production Monitoring & IoT

- Stream sensor data from industrial equipment.
- Enable predictive maintenance using anomaly detection.
- Real-time alerts to prevent downtime/defects.
- Use event pattern detection with plain SQL.

# LIMITATIONS & CONSIDERATIONS

- Requires upfront compute & storage setup
- Trade-offs: **latency** ↔ **throughput** ↔ **cost**
- Some **streaming complexities remain** (e.g., state mgmt)
- Best for **low-latency use cases**
- Ecosystem still maturing compared to Flink/Spark

## WHEN TO USE RISINGWAVE:

- You need real-time insights, not hourly/daily reports.
- You want to use SQL, not complex Java/Scala streaming code.
- You want a single system, not a complex stack.
- You need elastic scaling and cloud-native deployment.
- You want materialized views that update in real time.
- You want exactly-once processing guarantees.



thank you

